

A Review of Column-Oriented Datastores

By: Zach Pratt
Spring 2011

Introduction

- The problem: Internet-scale applications have additional requirements over and above those of lesser scale applications.
 - Reliably storing session state and application data for millions of users as apposed to thousands of users.
 - Handle high volumes of data by using dynamic partitioning
 - Replication for high availability
 - No single point of failure

RDBMS Limitations

- Traditional relational database systems aren't designed with the requirements of Internet-scale applications in mind.
 - Partitioning must be part of the initial logical data model, if its even supported.
 - Replicating to many nodes while maintaining strong consistency limits availability.
 - Single points of failure

Solution

- Use a non-relational datastore (NoSQL) that is designed with the following in mind:
 - Dynamic partitioning
 - Large scale replication
 - Totally decentralization
- There are many; I will cover a few column-oriented versions. See <http://nosql-database.org/> for a big list of other implementations.

Background

- There are many differences between relational databases and column-oriented datastores
 - Data models
 - CRUD operations
 - Support for ACID properties

Data Models

- Relational vs Column-Oriented
 - Relational:
 - DDL to define tables and constraints
 - Schema contained within the database
 - Rows of columns make up tables
 - Column-Oriented: Groupings of key-value pairs
 - Schema is primarily procedural
 - Groups of key-value pairs instead of tables

Example Data Model

BigTable and Cassandra

```
row_key: {  
  key_name1:key_value1  
  , ...  
  , key_nameN:key_valueN  
}
```

```
edu.wiu.appserver: {  
  architecture:x86_64  
  num_cpus:4  
  memory:16GB  
  ip_address:10.200.15.4  
}
```

Dynamo

```
key_name1:blob1  
...  
key_nameN:blobN
```

CRUD Operations

- Relational
 - DML for inserts, updates, deletes
 - SQL for queries
 - Both are fairly standard across RDBMS's.
- Column-Oriented
 - Each datastore has its own client-side API
 - `get()` for queries
 - `put()` for inserts and updates

ACID Properties

- ACID: **A**tomicity, **C**onsistency, **I**solation, **D**urability
- Atomicity: all or nothing transactions
- Consistency: ensure the data is valid
- Isolation: handling concurrent operations
- Durability: recover committed data after a failure

Theory

- Important column-oriented principles:
 - CAP Theorem
 - Eventual consistency
 - Consistent hashing

CAP Theorem

- Initially proposed by Eric Brewer
- Distributed systems of sufficient size cannot maintain:
 - Consistency
 - Availability
 - Partition tolerance
- Trade-offs must be made between consistency and availability in order to maintain high availability.


Eventual Consistency

- Strong consistency limits availability
 - Especially synchronous replication
- Eventual consistency
 - Data on all nodes will be consistent... eventually
 - Applications can specify their required level of consistency for reads and writes.

Consistent Hashing

- Used to support dynamic partitioning
 - Each node is assigned a token, which represents a range within the hash space for which it is responsible.
 - Keys are hashed and assigned to the node handling the range in which the hashed key falls under

Conceptual Hash Space



Implementations

- Google's BigTable
- Amazon's Dynamo
- Facebook's Cassandra

Google's BigTable

- Focuses on dynamically partitioning data
 - Does not focus on fault tolerance, therefore, replication is not a feature
- Data Model: Column Keys, Column Families, Rows, Timestamps
 - Semi-structured data
- Partitions data by row key range amongst a cluster of tablet servers.

Amazon's Dynamo

- Totally Decentralized
- Focuses on dynamic partitioning, large scale replication, fault tolerance
- Data Model: Single namespace for key-value pairs which results in an unstructured data model.
- Eventually consistent
- Has no security.

Facebook's Cassandra

- Combines parts of BigTable and Dynamo
- BigTable components:
 - Data Model: Column Keys, Column Families: Simple and Super, Rows
 - Durability: Memtables and SSTables with a commit log
- Dynamo components: eventual consistency, dynamic partitioning using consistent hashing.

Real World Use Case

- Storing IPTables log files from multiple servers
 - Log messages are often formatted as key-value pairs.
 - This solution would require a high write throughput.
 - Store and analyze large volumes of data

Use Case Data Model

```
192.168.1.2_Apr-23-18:33:12 {  
  FROMHOST-IP: 192.168.1.2  
  IN: eth0  
  OUT:  
  MAC: 00:00:00:00:00:00  
  SRC: 192.168.1.8  
  DST: 192.168.1.2  
  LEN: 200  
  TOS: 0x00  
  PREC: 0x00  
  TTL: 54  
  ID: 3000  
  PROTO: UDP  
  SPT: 53  
  DPT: 31232  
}
```

```
192.168.1.23_Apr-23-18:32:08 {  
  FROMHOST-IP: 192.168.1.23  
  IN: eth0  
  OUT:  
  MAC: 00:00:00:00:00:00  
  SRC: 192.168.1.8  
  DST: 192.168.1.23  
  LEN: 40  
  TOS: 0x00  
  PREC: 0x00  
  TTL: 32  
  ID: 561  
  PROTO: TCP  
  SPT: 8080  
  DPT: 30215  
  WINDOW: 65535  
  RES: 0x00  
  FLAG: SYN  
}
```

Conclusion

- Non-relational column-oriented systems have specific use cases.
- They are not meant to replace relational database systems
 - Some applications require transaction handling and strong consistency
 - Banking Applications
- They are often used together